



LEVEL 1 – PickAndInsertCube.Station1.merged.deltajoint_old (v2.1 LeRobot dataset): 136 episodes, 39 603 frames total, 20 Hz



each box = one `data/chunk-000/episode_XXXXXX.parquet` (+ matching videos/.../episode_XXXXXX.mp4 per camera)

LEVEL 2 – episode 009 zoomed: a sequence of 352 per-frame rows (this episode's actual length), 20 Hz $\Rightarrow$ 17.6 s



this dataset's `next.done` region is verified **CONTIGUOUS** from a “success onset” frame to the literal last frame in all 136 episodes (mean 41.2 frames marked True, out of a mean 291-frame episode) – see main text §1.

LEVEL 3 – one raw dataset row (frame 350), EXACT schema (verified: `meta/info.json` + `parquet` inspection)

Per-frame keys stored in the `parquet` (images live in the separate per-episode `.mp4` instead, referenced by `episode_index/frame_index`):

<code>action</code>	(7,) float32 – delta-joint target, this frame
<code>observation.state</code>	(7,) float32 – proprioceptive state at this frame
<code>observation.joint.pos_raw,</code> <code>.ee.pose_raw, .velocity, .effort,</code> <code>.acceleration</code>	(7,) float32 each – raw sensor channels (not used by the RL pipeline directly)
<code>observation.images.cam_left_wrist,</code> <code>.cam_right_wrist</code>	decoded from the episode's <code>.mp4</code> at load time
<code>next.done</code>	bool – True for frame 350 (inside the verified success-held region)
<code>next.reward</code>	float – 1.0 for frame 350 (identical mask to <code>next.done</code> , verified in all 136 episodes)
<code>timestamp, frame_index,</code> <code>episode_index, index, task_index</code>	bookkeeping (global vs. per-episode indexing)

This is what `.populate_offline_buffer()` reads per frame, BEFORE any n -step processing – see Figure “offline buffer population” for what happens next.